



# COMET What is next?


# Support for Emerging Dataflow Architectures

- Novel compiler technologies are fundamental to generate efficient computational graph for emerging dataflow architectures
- We are working on adding support to COMET to generate optimized code for emerging dataflow architectures

SambaNova SN10



Xilinx Versal



## Other Features

- **Sparse Transpose**
- Initial design to support code generation on **Xilinx Versal AI engine** and **Sambanova**
- Multiple frontends (**CometPy** and Rust **eDSL**)
- Integration with other tools:
  - **MCL** for execution on extremely heterogeneous systems
  - **Lamellar** for execution non distributed systems
- **HW/SW Co-Design**
  - COMET automatically inserts calls to novel hardware accelerators, real code of simulators (e.g., Aladdin)

## Conclusions

- Explosion of programming framework and heterogeneous architectures challenge traditional compiler development
- COMET is a compiler and support various frontend (NumPy and Rust) targeting computational chemistry and graph analytics domains
  - Based on MLIR
  - Support sparse and dense tensor algebra computation
- COMET employs high-level, domain specific optimizations and low-level, architecture specific optimization to generate efficient code for target architectures
- COMET targets various architecture, including parallel CPUs, GPUs, and FPGAs.
- Open source: <https://github.com/pnnl/COMET>



# References

- COMET: A domain-specific compilation of high-performance computational chemistry. Erdal Mutlu, Ruiqin Tian, Bin Ren, Sriram Krishnamoorthy, Roberto Gioiosa, Jacques Pienaar, Gokcen Kestor. International Workshop on Languages and Compilers for Parallel Computing. October 2020.
- A High Performance Sparse Tensor Algebra Compiler in MLIR. Ruiqin Tian, Luanzheng Guo, Jiajia Li, Bin Ren, Gokcen Kestor. 2021 IEEE/ACM 7th Workshop on the LLVM Compiler Infrastructure in HPC (LLVM-HPC). November, 2021.
- Union: A unified HW-SW co-design ecosystem in MLIR for evaluating tensor operations on spatial accelerators. Geonhwa Jeong, Gokcen Kestor, Prasanth Chatarasi, Angshuman Parashar, Po-An Tsai, Sivasankaran Rajamanickam, Roberto Gioiosa, Tushar Krishna. 2021 30th International Conference on Parallel Architectures and Compilation Techniques (PACT). September 2021.
- ReACT: Redundancy-Aware Code Generation for Tensor Expressions. Tong Zhou, Ruiqin Tian, Rizwan A Ashraf, Roberto Gioiosa, Gokcen Kestor, Vivek Sarkar. 2022 31st International Conference on Parallel Architectures and Compilation Techniques (PACT). October 2022
- Guo L., R.A. Ashraf, R.D. Friese, and G. Kestor. 10/10/2022. "Towards Supporting Semiring in MLIR-based COMET Compiler." Presented by G. Kestor at the 31st International Conference on Parallel Architectures and Compilation Techniques (PACT), Chicago, Illinois.





- # Thank you

The diagram illustrates a distributed system architecture. It features four light blue cubes, each with a green square on its front face, representing nodes or servers. These nodes are arranged in a square pattern and are interconnected by a network of lines. In the center of the network is a circular hub containing a stylized pyramid with horizontal stripes in blue, green, and yellow. The entire system is set against a light blue background with a subtle grid pattern.

6